

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense
import random

2026-02-05 10:23:30.638649: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to
register cuFFT factory: Attempting to register factory for plugin
cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are
written to STDERR
E0000 00:00:1770287010.814596      55 cuda_dnn.cc:8579] Unable to
register cuDNN factory: Attempting to register factory for plugin
cuDNN when one has already been registered
E0000 00:00:1770287010.867149      55 cuda_blas.cc:1407] Unable to
register cuBLAS factory: Attempting to register factory for plugin
cuBLAS when one has already been registered
W0000 00:00:1770287011.273873      55 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.
W0000 00:00:1770287011.273908      55 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.
W0000 00:00:1770287011.273911      55 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.
W0000 00:00:1770287011.273913      55 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.

data = """
artificial intelligence is transforming modern society.
it is used in healthcare finance education and transportation.
machine learning allows systems to improve automatically with
experience.
data plays a critical role in training intelligent systems.
large datasets help models learn complex patterns.
deep learning uses multi layer neural networks.
neural networks are inspired by biological neurons.
each neuron processes input and produces an output.
training a neural network requires optimization techniques.
gradient descent minimizes the loss function.

natural language processing helps computers understand human language.
text generation is a key task in nlp.
language models predict the next word or character.
recurrent neural networks handle sequential data.
```

```
lstm and gru models address long term dependency problems.  
transformer models changed the field of nlp.  
they rely on self attention mechanisms.  
attention allows the model to focus on relevant context.  
  
education is being improved using artificial intelligence.  
intelligent tutoring systems personalize learning.  
  
ethical considerations are important in artificial intelligence.  
ai systems should be designed responsibly.  
  
text generation models can create stories poems and articles.  
generated text should be meaningful and coherent.  
  
continuous learning is essential in the field of ai.  
programming skills are important for ai engineers.  
"""
```

N-Gram Model

```
words = data.lower().replace("\n", " ").split()  
  
ngram_model = {}  
  
for i in range(len(words)-1):  
    w1 = words[i]  
    w2 = words[i+1]  
  
    if w1 not in ngram_model:  
        ngram_model[w1] = []  
  
    ngram_model[w1].append(w2)  
  
print("Total Keys in Ngram Model:", len(ngram_model))  
  
Total Keys in Ngram Model: 139  
  
def generate_ngram_text(seed, n_words=20):  
    result = [seed]  
  
    for _ in range(n_words):  
        last = result[-1]  
        if last in ngram_model:  
            next_word = random.choice(ngram_model[last])  
            result.append(next_word)  
        else:  
            break  
  
    return " ".join(result)
```

```

print(generate_ngram_text("artificial", 25))

artificial intelligence. intelligent tutoring systems personalize
learning. ethical considerations are inspired by biological neurons.
each neuron processes input and coherent. continuous learning allows
systems personalize learning.

```

LIMITATIONS OF NGRAM MODEL:

- Only looks at previous one word (short memory)
- Cannot understand long context
- Generated text may repeat or lose meaning
- No deep learning or semantic understanding

RNN Model:

```

tokenizer = Tokenizer()
tokenizer.fit_on_texts([data])

total_words = len(tokenizer.word_index) + 1
print("Total Vocabulary:", total_words)

Total Vocabulary: 134

input_sequences = []

for line in data.split("\n"):
    token_list = tokenizer.texts_to_sequences([line])[0]

    for i in range(1, len(token_list)):
        n_gram_seq = token_list[:i+1]
        input_sequences.append(n_gram_seq)

max_len = max(len(x) for x in input_sequences)

input_sequences = pad_sequences(input_sequences, maxlen=max_len,
padding='pre')

X = input_sequences[:, :-1]
y = input_sequences[:, -1]

y = tf.keras.utils.to_categorical(y, num_classes=total_words)

print(X.shape, y.shape)

(167, 8) (167, 134)

model = Sequential([
    Embedding(input_dim=total_words,
              output_dim=64,

```

```

        input_shape=(max_len-1,)),
        SimpleRNN(128),
        Dense(total_words, activation='softmax')
    ])

model.compile(loss='categorical_crossentropy',
               optimizer='adam',
               metrics=['accuracy'])

model.summary()

```

/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/embedding.py:100: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

super().__init__(**kwargs)

I0000 00:00:1770287033.336214 55 gpu_device.cc:2019] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 13757 MB memory: -> device: 0, name: Tesla T4, pci bus id: 0000:00:04.0, compute capability: 7.5

I0000 00:00:1770287033.342162 55 gpu_device.cc:2019] Created device /job:localhost/replica:0/task:0/device:GPU:1 with 13757 MB memory: -> device: 1, name: Tesla T4, pci bus id: 0000:00:05.0, compute capability: 7.5

Model: "sequential"

Layer (type)	Output Shape
Param #	
embedding (Embedding)	(None, 8, 64)
8,576	
simple_rnn (SimpleRNN)	(None, 128)
24,704	
dense (Dense)	(None, 134)
17,286	

Total params: 50,566 (197.52 KB)

Trainable params: 50,566 (197.52 KB)

```
Non-trainable params: 0 (0.00 B)

history = model.fit(X, y, epochs=100, verbose=1)

Epoch 1/100

WARNING: All log messages before absl::InitializeLog() is called are
written to STDERR
I0000 00:00:1770287039.782705      125 service.cc:152] XLA service
0x7a4680004580 initialized for platform CUDA (this does not guarantee
that XLA will be used). Devices:
I0000 00:00:1770287039.782748      125 service.cc:160] StreamExecutor
device (0): Tesla T4, Compute Capability 7.5
I0000 00:00:1770287039.782754      125 service.cc:160] StreamExecutor
device (1): Tesla T4, Compute Capability 7.5
I0000 00:00:1770287040.098936      125 cuda_dnn.cc:529] Loaded cuDNN
version 91002

1/6 ━━━━━━━━ 14s 3s/step - accuracy: 0.0312 - loss: 4.8933

I0000 00:00:1770287041.321151      125 device_compiler.h:188] Compiled
cluster using XLA! This line is logged at most once for the lifetime
of the process.

6/6 ━━━━━━━━ 4s 210ms/step - accuracy: 0.0156 - loss:
4.9022
Epoch 2/100
6/6 ━━━━━━━━ 0s 7ms/step - accuracy: 0.0796 - loss: 4.8116

Epoch 3/100
6/6 ━━━━━━━━ 0s 7ms/step - accuracy: 0.1443 - loss: 4.7097

Epoch 4/100
6/6 ━━━━━━━━ 0s 7ms/step - accuracy: 0.1289 - loss: 4.5894

Epoch 5/100
6/6 ━━━━━━━━ 0s 7ms/step - accuracy: 0.1014 - loss: 4.5135

Epoch 6/100
6/6 ━━━━━━━━ 0s 7ms/step - accuracy: 0.1074 - loss: 4.4213

Epoch 7/100
6/6 ━━━━━━━━ 0s 7ms/step - accuracy: 0.1563 - loss: 4.2166

Epoch 8/100
6/6 ━━━━━━━━ 0s 7ms/step - accuracy: 0.2265 - loss: 4.1719

Epoch 9/100
6/6 ━━━━━━━━ 0s 7ms/step - accuracy: 0.1721 - loss: 4.0914

Epoch 10/100
```

```
6/6 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.1611 - loss: 4.0109
Epoch 11/100
6/6 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.1530 - loss: 3.8803
Epoch 12/100
6/6 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.2120 - loss: 3.8193
Epoch 13/100
6/6 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.2337 - loss: 3.6625
Epoch 14/100
6/6 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.2559 - loss: 3.4858
Epoch 15/100
6/6 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.3259 - loss: 3.3753
Epoch 16/100
6/6 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.3962 - loss: 3.2449
Epoch 17/100
6/6 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.4253 - loss: 3.0899
Epoch 18/100
6/6 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.4104 - loss: 3.0397
Epoch 19/100
6/6 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.4872 - loss: 2.8808
Epoch 20/100
6/6 ━━━━━━━━━━ 0s 7ms/step - accuracy: 0.5324 - loss: 2.6109
Epoch 21/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.6622 - loss: 2.4946
Epoch 22/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.6524 - loss: 2.4148
Epoch 23/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.6601 - loss: 2.3209
Epoch 24/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.6841 - loss: 2.1140
Epoch 25/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.7404 - loss: 1.9837
Epoch 26/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.6835 - loss: 1.9439
Epoch 27/100
```

```
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.7345 - loss: 1.8326
Epoch 28/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.7758 - loss: 1.6328
Epoch 29/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.7232 - loss: 1.5600
Epoch 30/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.7819 - loss: 1.5176
Epoch 31/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.8080 - loss: 1.3870
Epoch 32/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.8325 - loss: 1.3174
Epoch 33/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.8240 - loss: 1.2960
Epoch 34/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.8647 - loss: 1.1160
Epoch 35/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.8520 - loss: 1.1602
Epoch 36/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.8381 - loss: 1.1263
Epoch 37/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.8342 - loss: 1.0644
Epoch 38/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.8181 - loss: 0.9858
Epoch 39/100
6/6 ━━━━━━━━━━ 0s 9ms/step - accuracy: 0.8610 - loss: 0.9850
Epoch 40/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.8923 - loss: 0.8999
Epoch 41/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.8997 - loss: 0.8410
Epoch 42/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9179 - loss: 0.8402
Epoch 43/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9045 - loss: 0.8198
Epoch 44/100
```

```
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9208 - loss: 0.7731
Epoch 45/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9393 - loss: 0.7140
Epoch 46/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9349 - loss: 0.6496
Epoch 47/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9144 - loss: 0.6779
Epoch 48/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9360 - loss: 0.5800
Epoch 49/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9364 - loss: 0.6196
Epoch 50/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9289 - loss: 0.5921
Epoch 51/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9688 - loss: 0.4813
Epoch 52/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9715 - loss: 0.4599
Epoch 53/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9680 - loss: 0.4880
Epoch 54/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9689 - loss: 0.4904
Epoch 55/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9695 - loss: 0.4127
Epoch 56/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9743 - loss: 0.4008
Epoch 57/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9797 - loss: 0.4314
Epoch 58/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9702 - loss: 0.3791
Epoch 59/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9754 - loss: 0.3400
Epoch 60/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9827 - loss: 0.3708
Epoch 61/100
```

```
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9769 - loss: 0.3915
Epoch 62/100
6/6 ━━━━━━━━ 0s 10ms/step - accuracy: 0.9814 - loss: 0.3190
Epoch 63/100
6/6 ━━━━━━ 0s 8ms/step - accuracy: 0.9688 - loss: 0.3395
Epoch 64/100
6/6 ━━━━ 0s 8ms/step - accuracy: 0.9650 - loss: 0.3371
Epoch 65/100
6/6 ━━ 0s 9ms/step - accuracy: 0.9717 - loss: 0.3010
Epoch 66/100
6/6 ━ 0s 8ms/step - accuracy: 0.9807 - loss: 0.3296
Epoch 67/100
6/6 0s 8ms/step - accuracy: 0.9922 - loss: 0.2604
Epoch 68/100
6/6 0s 8ms/step - accuracy: 0.9922 - loss: 0.2605
Epoch 69/100
6/6 0s 8ms/step - accuracy: 0.9801 - loss: 0.2277
Epoch 70/100
6/6 0s 8ms/step - accuracy: 0.9900 - loss: 0.2559
Epoch 71/100
6/6 0s 8ms/step - accuracy: 0.9710 - loss: 0.2247
Epoch 72/100
6/6 0s 8ms/step - accuracy: 0.9807 - loss: 0.2165
Epoch 73/100
6/6 0s 8ms/step - accuracy: 0.9881 - loss: 0.2274
Epoch 74/100
6/6 0s 8ms/step - accuracy: 0.9963 - loss: 0.1817
Epoch 75/100
6/6 0s 8ms/step - accuracy: 0.9881 - loss: 0.1913
Epoch 76/100
6/6 0s 8ms/step - accuracy: 0.9881 - loss: 0.1840
Epoch 77/100
6/6 0s 8ms/step - accuracy: 0.9974 - loss: 0.1836
Epoch 78/100
```

```
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9926 - loss: 0.1664
Epoch 79/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9948 - loss: 0.1709
Epoch 80/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9926 - loss: 0.1756
Epoch 81/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9926 - loss: 0.1287
Epoch 82/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9963 - loss: 0.1424
Epoch 83/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9963 - loss: 0.1340
Epoch 84/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9948 - loss: 0.1374
Epoch 85/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9974 - loss: 0.1348
Epoch 86/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9974 - loss: 0.1199
Epoch 87/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9888 - loss: 0.1112
Epoch 88/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9948 - loss: 0.1286
Epoch 89/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9948 - loss: 0.1073
Epoch 90/100
6/6 ━━━━━━━━━━ 0s 9ms/step - accuracy: 0.9881 - loss: 0.1084
Epoch 91/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9807 - loss: 0.1015
Epoch 92/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9829 - loss: 0.1066
Epoch 93/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9881 - loss: 0.1056
Epoch 94/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9881 - loss: 0.0964
Epoch 95/100
```

```
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9974 - loss: 0.0829
Epoch 96/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9926 - loss: 0.0937
Epoch 97/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9963 - loss: 0.0844
Epoch 98/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9855 - loss: 0.0884
Epoch 99/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9983 - loss: 0.0789
Epoch 100/100
6/6 ━━━━━━━━━━ 0s 8ms/step - accuracy: 0.9926 - loss: 0.0775

def generate_rnn_text(seed_text, next_words=20):
    for _ in range(next_words):

        token_list = tokenizer.texts_to_sequences([seed_text])[0]
        token_list = pad_sequences([token_list], maxlen=max_len-1,
padding='pre')

        predicted = model.predict(token_list, verbose=0)
        predicted_word = tokenizer.index_word[np.argmax(predicted)]

        seed_text += " " + predicted_word

    return seed_text

print(generate_rnn_text("artificial intelligence", 25))

artificial intelligence is transforming modern society and of ai
systems should be designed responsibly coherent an output term loss
problems human help experience context computers understand human

print("\n====")
print(" N-GRAM vs RNN COMPARISON")
print("=====\\n")

seed = "artificial intelligence"

print("Seed Text:", seed)
print("\n--- NGRAM OUTPUT ---")
print(generate_ngram_text("artificial", 25))

print("\n--- RNN OUTPUT ---")
print(generate_rnn_text(seed, 25))

print("\n--- OBSERVATION ---")
```

```

print("NGRAM:")
print("- Local word prediction only")
print("- May break sentence flow")

print("\nRNN:")
print("- Understands sequence patterns")
print("- More meaningful and coherent text")

=====
N-GRAM vs RNN COMPARISON
=====

Seed Text: artificial intelligence

--- NGRAM OUTPUT ---
artificial intelligence. ai engineers.

--- RNN OUTPUT ---
artificial intelligence is transforming modern society and of ai
systems should be designed responsibly coherent an output term loss
problems human help experience context computers understand human

--- OBSERVATION ---
NGRAM:
- Local word prediction only
- May break sentence flow

RNN:
- Understands sequence patterns
- More meaningful and coherent text

import gradio as gr

def rnn_ui(seed_text, length):
    return generate_rnn_text(seed_text, int(length))

demo = gr.Interface(
    fn=rnn_ui,
    inputs=[
        gr.Textbox(label="Enter Seed Text", value="artificial
intelligence"),
        gr.Slider(5, 50, value=20, step=1, label="Generate Words")
    ],
    outputs=gr.Textbox(label="Generated Text"),
    title="RNN Text Generation UI",
    description="GenAI Lab-4 – Text Generation using RNN"
)

demo.launch(debug=True)

```

```
* Running on local URL: http://127.0.0.1:7860
It looks like you are running Gradio on a hosted Jupyter notebook,
which requires `share=True`. Automatically setting `share=True` (you
can turn this off by setting `share=False` in `launch()` explicitly).
```

```
* Running on public URL: https://1f6e629edaf5aca289.gradio.live
```

```
This share link expires in 1 week. For free permanent hosting and GPU
upgrades, run `gradio deploy` from the terminal in the working
directory to deploy to Hugging Face Spaces
(https://huggingface.co/spaces)
```

```
<IPython.core.display.HTML object>
```

```
Created dataset file at: .gradio/flagged/dataset1.csv
```